

Commands and Function Learning

-----P_1-----

Commands:

echo :

- Print a message to the terminal

- Display the value of variables ex -

 - \$name="Shikhar"

 - \$echo "My

 - name is \$name"

- Formatting output (with options)

 - e ? Enable interpretation

- of backslash escapes

 - \$echo -e "Line1\nLine2\tTabbed"

- Redirect output to

- a file

 - \$echo "Hello - original file (symlink demo)" > original.txt

ln : [Important]

1.Hard Link (default)

\$ln original.txt hardlink.txt

- Creates another

- name (hardlink.txt) pointing to the same inode as original.txt.

- Both share the same data

blocks.

- If one is modified, the other reflects changes.

- If the original is deleted,

- the data still exists (through the hard link).

2.Symbolic Link (Soft Link)

\$ln -s original.txt

symlink.txt

- Creates a shortcut (path reference) to the original file.

- If the original

- file is deleted, the symlink becomes broken.

- Can link to directories and across

- filesystems.

- s ? Create symbolic link (soft link).

ls (List File) : used to list file

- Simple list :\$ls

 - Long listing :\$ls -l

 - Hidden files included : \$ls -a

- Human-readable sizes : \$ls -lh

 - Recursive listing: \$ls -R

 - List with inode number :

\$ls -li

cat (concatenate) : used to view, create, append, merge, and copy files

- Display

- contents of a file : \$cat file.txt

- View multiple files together : \$cat file1.txt

- file2.txt

- Create a new file : \$cat > newfile.txt

- Append content to a file : \$cat

- >> existing.txt

- Combine multiple files into one : \$cat file1.txt file2.txt >

- merged.txt

- Number the lines in output : \$cat -n file.txt

- Quickly copy a file : \$cat

- file.txt > copy.txt

chmod (change mode) :

- \$chmod 755 script.sh

- > Owner (7),

- Group (5), Others (5)

```
mkfifo (Make Named Pipe) :
    -creates a named pipe FIFO = First In,
First Out
    -mkfifo creates a named pipe FIFO = First In, First Out.
    -It is a special
file that processes can use for inter-process communication IPC.
    -Unlike normal pipes (|),
a named pipe exists as a file in the filesystem.
=>$mkfifo mypipe
=>$ls -l
```

```
prw-r--r-- 1 shikhar users 0 Sep  2 23:45 mypipe
p at the start (prw-...) shows it's a
pipe.
```

Permission Codes:

```
Read = 4
Write = 2
Execute = 1
rwx = 4+2+1 = 7
rw- = 4+2+0 = 6
r-- =
4+0+0 = 4
```

```
----- C Programming : -----
```

```
open(orig,
O_WRONLY | O_CREAT | O_TRUNC, 0644); --> Opens or creates file descriptor for
writing.
close(); --> close file descriptor
```

```
write(); --> Writes it into file.
read();
--> Read data into var from file.
```

```
symlink(orig, slink); --> Creates a symbolic link
named SoftLink named hplac.txt pointing to file.
link(orig, hlink); --> Creates a hard
link
unlink(orig); --> Deletes the directory entry for
original_hard.txt.
mkfifo("myfifo", 0666); --> tries to create a named pipe with
read/write permissions for all.
fork(); --> duplicates the process.
```

```
perror(); --> used
to print a human-readable error message to stderr standard error.
```

```
lstat(slink, &st_link);
--> gets info about the symlink itself.
stat(orig, &st1); --> function fills st1 with
metadata about the file named by orig.
```

```
readlink(); --> reads the contents of a symlink the
path it points
to.
```

```
-----P_2-----
```

```
while(1) --> infinite
loop
```

```
-----P_3-----
```

```
----- C Programming : -----
```

```
creat("myfile.txt", 0644); --> Calls creat to create or
truncate the file 'myfile.txt' and open it for writing. The mode 0644 is an octal
```

literal.

-----P_4-----

----- C Programming :

O_EXCL(Exclusive creation) ensures that the open fails if the file already exists.

open("myfile.txt", O_RDWR | O_CREAT | O_EXCL, 0644); -->
O_EXCL only meaningful with O_CREAT ? fail if the file already exists.

-----P_5-----

Commands:

after compilation run this
\$./create_files

& --> runs the compiled program create_files as a background process, letting you continue using the terminal while it runs.

Still tied to your terminal

session.

o/p --> [1] 40877 === job number (like [1]) and the PID (process ID)

\$nohup ./create_files & --> nohup = no hangup.

Tells the system to ignore SIGHUP signals, so the program keeps running even after logout or terminal close.

----- C Programming :

snprintf(filename, sizeof(filename), "file%d.txt", i + 1);

--> snprintf ensures the string fits in filename safe version of sprintf.

(Used to assign filename in char

array.)

-----P_6-----

----- C Programming :

1. read(int fd, void *buf, size_t count):

=>Purpose:

->read() is a low-level system call that reads raw bytes from a file descriptor into memory.

->Works for files, pipes, sockets, devices – basically anything that can be represented by a file descriptor (fd).

=>Parameters

->fd

? File descriptor

-Obtained from open(), pipe(), socket(), etc.

- 0 = standard input (stdin)

- 1 = standard output (stdout)

- 2 = standard error (stderr)

->buf ? Pointer to a buffer

(array in memory) where data will be stored.

->count ? Maximum number of bytes

to read.

=> Return Value

-> > 0 ? Number of bytes actually

read.

-> 0 ? End of file (EOF).

-> -1 ? Error (sets errno).

```

// Calls the low-level read() system call:
// 0 is the file descriptor for STDIN.
//
buffer is where bytes will be stored.
// sizeof(buffer) = 100 is the maximum number of
bytes to read.
// read() blocks until input is available, EOF is reached, or an error
occurs.
// Return value:
-->n > 0 ? number of bytes actually read (?
100).
-->n == 0 ? EOF (no more input).
-->n == -1 ? error occurred
(and errno is set).

```

2. write(int fd, const void *buf, size_t count):
=>Purpose:

->write() is a low-level system call that writes raw bytes from memory (buf) to a file descriptor (fd).

->Works for files, pipes, sockets, devices – basically anything that can be represented by a file descriptor.

=>Parameters:

1. fd ? File
descriptor

Obtained from open(), pipe(), socket(), etc.
0 = standard

input (stdin)

1 = standard output (stdout)
2 = standard error (stderr)

2. buf ? Pointer to a buffer (array in memory) where data will be stored.

3.

count ? Maximum number of bytes to read.

4. Return Value

> 0 ?

Number of bytes actually read.

0 ? End of file (EOF).

-1 ? Error

(sets
errno).

-----P_7-----

Commands:

copy <filename> <destination>

----- C Programming : -----

1.main(int argc, char
*argv[]) :

=> Parameters:

1.argc ? Argument Count

An integer

representing the number of command-line arguments passed to the program.

Always ? 1

because the program name itself counts as the first argument.

2.argv ? Argument

Vector

An array of strings (char pointers), holding the arguments.

argv[0]

? name of the program (string).

argv[1] ... argv[argc-1] ? actual arguments

passed by the user.

argv[argc] ? guaranteed to be NULL (marks the end).

ex -

./mycp <source> <destination>

2.while ((n = read(fd1, buffer, SIZE)) > 0)

{

if (write(fd2, buffer, n) != n) {}

```
}
```

-----P_8-----

-----C Programming :-----

```
Read file line by line
while ((n = read(fd, buffer, SIZE)) >
0) { // Read Line by Line
    for (int i = 0; i < n; i++) {
        line[idx++] =
buffer[i];
        if (buffer[i] == '\n') { // found a line
write(1, line, idx); // print to STDOUT
            idx = 0; //
reset line buffer
        }
    }
}
```

-----P_9-----

-----C Programming :-----

```
fprintf(stderr, "Usage: %s <filename>\n",
argv[0]) :
    ->It is just like printf(), but instead of printing to stdout
(screen), it prints (writes formatted text) to a file stream.
    ->You control
where the output goes (file, terminal, socket, etc.).
```

```
struct stat fileStat
```

->used to define struct of file.

```
stat(filename, &fileStat):
```

->fills fileStat with metadata about the file argv[1].

```
argv[1] : Prints
the name of the file being examined.
(long)fileStat.st_ino : Prints inode number
(unique identifier for the file on disk).
(long)fileStat.st_nlink : Number of hard
links ? how many directory entries (names) point to this file.
fileStat.st_uid : UID
(User ID) of file's owner.
fileStat.st_gid : GID (Group ID) of file's group.

(long)fileStat.st_size : File size in bytes (actual data length).

(long)fileStat.st_blksize : Preferred block size for filesystem I/O (not file size). Helps
optimize read()/write()
(long)fileStat.st_blocks : Number of disk blocks allocated to
the file.
ctime(&fileStat.st_atime) : when the file was last read.

ctime(&fileStat.st_mtime) : when file content last changed.

ctime(&fileStat.st_ctime) : when metadata (permissions, ownership, links) last
changed.
```

-----P_10-----

Commands :

od :

Octal Dump

(but it can also show hex, chars, etc.).

It shows the raw contents of a file in

a readable way.

-Default: octal representation.

-With

-c: show characters (printable ASCII or \0, \n, etc.).

-With -x: show

hexadecimal.

? Example:

echo "ABC" > f.txt

od -c f.txt

----- C Programming :

open("holefile.txt", O_RDWR | O_CREAT | O_TRUNC, 0644)

O_TRUNC :

->If the file already exists ? its contents are erased

(truncated to length 0).

->If the file does not exist ? nothing to truncate,

but since O_CREAT is also specified, a new empty file will be created.

->File

permissions come from the 3rd argument (0644 here) when creating a new file.

lseek(fd,

10, SEEK_CUR); lseek return value (new offset)

// move file pointer forward by 10 bytes

(creating a hole)

lseek(int fd, off_t offset, int whence);

=>Purpose

-Moves (or queries) the file offset (also called the "file pointer").

-The file offset tells the kernel where in the file the next read() or write() will happen.

=>Parameters

-fd ? file descriptor (from open()).

-offset ?

number of bytes to move.

-whence ? starting position, one of:

-SEEK_SET ? from beginning of file.

-SEEK_CUR ? from current file offset.

-SEEK_END ? from end of

file.

-----P_11-----

----- C Programming :

dup(fd) :

->dup(oldfd) creates a new file

descriptor that refers to the same open file description as oldfd. It returns the

lowest-numbered unused

file descriptor on success, or -1 on error.

dup2(fd, 10) :

-> If fd is invalid ? returns -1 with errno = EBADF.

-> If fd == 10 ? nothing changes, return 10 immediately.

-> If 10 is

already open ? it is closed first (to avoid resource leaks).

-> 10 is then

made a duplicate of fd.

-Now fd and 10 share the same open file description

(file offset, O_APPEND, etc.).

-Writes/reads on one affect the other.

fcntl(fd, F_DUPFD, 20) :

-> Duplicates fd into a new file descriptor.

-> The new file descriptor will be greater than or equal to start_fd (20 in this case).

-> Kernel finds the lowest available fd ? 20 and assigns it.

-> Both descriptors (fd and newfd) point to the same open file description (same offset, same flags).

-----P_12-----

----- C Programming : -----

- fcntl(): is a Swiss army knife for file descriptors.

- It can

duplicate, modify, query, or lock file descriptors.

- It is often used in:

-

Non-blocking I/O

- File locking

- Descriptor duplication (like in shells,

pipes, redirection)

- Signal-driven I/O

=> flags = fcntl(fd, F_GETFL);

->Get file status flags using fcntl

->Duplicating File Descriptors

int newfd = fcntl(fd, F_DUPFD, 20);

->File Descriptor Flags

F_GETFD ? Get per-FD flags (like FD_CLOEXEC).

F_SETFD ? Set per-FD flags.

Example:

fcntl(fd, F_SETFD, FD_CLOEXEC);

->File Status Flags

F_GETFL ? Get file status flags (set during open()).

Includes access mode

(O_RDONLY, O_WRONLY, O_RDWR) and status flags (O_APPEND, O_NONBLOCK, O_SYNC, etc.).

F_SETFL ? Change file status flags.

int flags = fcntl(fd, F_GETFL);

fcntl(fd, F_SETFL, flags | O_NONBLOCK)

Summary of Common Commands :

Command	Action
`F_DUPFD`	Duplicate fd
? given number	
`F_DUPFD_CLOEXEC`	Duplicate fd with
`FD_CLOEXEC`	
`F_GETFD`	Get fd flags (like
`FD_CLOEXEC`)	
`F_SETFD`	Set fd flags
`F_GETFL`	Get file status flags (`O_APPEND`,
`O_NONBLOCK`, etc.)	

	`F_SETFL`	Set file status flags	
	`F_GETLK`	Get lock info	
	`F_SETLK`	Set lock (non-blocking)	
	`F_SETLKW`	Set lock (blocking)	
	`F_SETOWN`	Set	
owner process for async I/O			
	`F_GETOWN`	Get owner	
process for async I/O			

=> accmode = flags & O_ACCMODE :

-> // Extract access mode (O_RDONLY / O_WRONLY / O_RDWR)
 -> O_ACCMODE is a
 mask defined in <fcntl.h>.
 -> flags & O_ACCMODE extracts only the access
 mode bits (ignores other flags like O_APPEND).

accmode == O_RDONLY --> "Read
 only\n"
 accmode == O_WRONLY --> "Write only\n"
 accmode == O_RDWR
 --> "Read/Write\n"

flags & O_APPEND --> "O_APPEND flag is
 set\n"
 flags & O_NONBLOCK --> "O_NONBLOCK flag is set\n"
 flags
 & O_SYNC --> "O_SYNC flag is set\n"
 flags & O_CREAT -->
 "O_CREAT flag is set (only relevant at open
 time)\n"

-----P_13-----

----- C Programming :

fd_set readfds; ----> fd_set is implemented as a bit
 array.
 ->fd_set readfds; - a bitmap structure used by select() to represent a
 set of file descriptors to monitor for readability.
 struct timeval tv;

->struct timeval tv; - timeout structure with seconds (tv_sec) and microseconds
 (tv_usec).

FD_ZERO(&readfds); ----> clears the set of file descriptors so you can
 safely add the ones you want to monitor with select().
 -> This initializes
 (clears) the fd_set structure.
 -> It sets all bits in readfds to 0, meaning no
 file descriptors are being monitored initially.
 -> clears the fd set
 (initializes to empty).

FD_SET(STDIN_FILENO, &readfds); ----> sets the bit
 corresponding to fd to 1 && add fd to monitor.
 -> adds STDIN (file
 descriptor 0) to the set of descriptors select() should watch for readability.

-> After this, readfds represents the set {0}.
 -> FD_SET(fd, &fdset) is a
 macro defined in <sys/select.h>.
 -> It adds the file descriptor fd to the
 file descriptor set fdset.
 -> In your case:
 - STDIN_FILENO ?
 this is 0 (standard input).
 - readfds ? the set of FDs you want select() to

monitor for readability.

-> So this line means:

"Monitor

standard input (keyboard input) for readability in the select() system call."

```
tv.tv_sec
```

```
= 10;
```

```
tv.tv_usec = 0;
```

->Sets the timeout to 10.0 seconds. tv_sec = whole

seconds, tv_usec = microseconds.

->Note: On many systems select() may modify tv

to reflect remaining time – so if you plan to reuse tv in a loop, reset it before each select() call.

```
fflush(stdout);
```

->ensures the message appears before

select() blocks (stdout is often line-buffered only when connected to terminal, but flushing is a safe practice).

```
retval = select(STDIN_FILENO + 1, &readfds, NULL, NULL,
&tv);
```

->select(nfds, &readfds, NULL, NULL, &tv) waits until one of

the monitored file descriptors is ready for reading, or until the timeout expires.

```
FD_ISSET(STDIN_FILENO, &readfds):
```

->tests if STDIN was marked ready by

```
select().
```

read(STDIN_FILENO, ...) to actually read the available bytes.

```
read()
```

returns:

- n > 0 = number of bytes read,

- n == 0 = EOF (e.g., user pressed Ctrl-D or input closed),

- n == -1 =

error.

->select() modifies its fd_set and timeval arguments. If you want to call select() again, you must reinitialize readfds and tv before each call.

->EOF behavior:

If the user sends EOF (Ctrl-D on an empty line), select() will report STDIN as readable and read() will return 0 (EOF).

Your code currently ignores n == 0 case, so nothing will be printed; you might want to handle EOF explicitly.

->Signals: If a signal interrupts

select() it returns -1 and errno == EINTR. Many programs retry select() in that case.

->Atomicity & line discipline: When STDIN is a terminal, input is typically line-buffered – the kernel won't mark STDIN readable until the user hits Enter.

->Alternatives: For large numbers of descriptors or high-performance servers, poll(), epoll (Linux) or kqueue (BSD/macOS) are often preferred.

-----P_14-----

----- C Programming : -----

```
fprintf()
```

->used to print formatted output into a file (or stream).

->Similar to printf(), but instead of always writing to stdout (screen),

you can write to any file or stream (FILE *).

printf() --> always prints to stdout.

sprintf() --> writes formatted text into a string buffer.

fprintf() -->

writes formatted text into a file or stream.

```

        S_ISREG(st.st_mode)  -->
Type: Regular file\n"
        S_ISDIR(st.st_mode)  -->      Type:
Directory\n"
        S_ISCHR(st.st_mode)  -->      Type: Character device\n"

        S_ISBLK(st.st_mode)  -->      Type: Block device\n"

S_ISFIFO(st.st_mode)  -->      Type: FIFO/pipe\n"
        S_ISLNK(st.st_mode)
-->      Type: Symbolic link\n"
        S_ISSOCK(st.st_mode)  -->      Type:
Socket\n"
        else  -->      Type:
Unknown\n"

```

-----P_15-----

----- C Programming :

```

extern char **environ;
    ->environ is a global variable
defined by libc

```

```

char **env = environ;
while (*env) {
    printf("%s\n",
*env);
    env++;
} :
    -> loop continues while *env is not NULL.

```

-> Prints the current environment string (format: KEY=VALUE) followed by a newline.

-----P_16-----

----- C Programming :

```

===== reader.c
=====

```

```

struct flock lock;
    ->Defines lock of type struct
flock, which describes a file lock (type, region, etc.).

```

```

Fill the struct flock for
READING :
    - l_type = F_RDLCK --> request a read lock (shared lock: multiple readers
allowed, but no writers).
    - l_whence = SEEK_SET --> offset is relative to the start
of the file.
    - l_start = 0 --> lock starts at byte 0.
    - l_len = 0 -->
lock to the end of the file (special meaning: "lock entire file").

```

```

fcntl(fd,
F_SETLKW, &lock) == -1 :
    - fcntl() with command F_SETLKW = "Set Lock,
Wait":
    - Tries to apply the lock (lock struct).
    - If another process holds
a conflicting lock (e.g., a write lock), this call blocks until it becomes available.

```

```

getchar();

```

```

lock.l_type = F_UNLCK;
fcntl(fd, F_SETLK, &lock);:

```

- Changes the lock type to F_UNLCK (unlock).
- Calls fcntl() with F_SETLK (non-blocking) to release the lock.

```
===== writer.c =====
```

Fill in the fields of the struct flock:

- l_type = F_WRLCK --> request a write lock (exclusive). Only one writer is allowed; no other readers or writers can hold locks at the same time.
- l_whence = SEEK_SET --> lock start is relative to the beginning of the file.
- l_start = 0 --> start from byte 0.
- l_len = 0 --> lock until the end of the file (special meaning = "entire file").

```
fcntl(fd, F_SETLKW,
&lock) == -1
    - F_SETLKW = Set Lock and Wait.
    - If another process already
holds a lock that conflicts (e.g., another write lock or a read lock), this call will block
until the lock can be acquired.
    - If fcntl fails, print error, close the file, and
exit.
```

```
lock.l_type = F_UNLCK;    // Unlock
fcntl(fd, F_SETLK, &lock);
-
Changes l_type to F_UNLCK ? unlock.
- Calls fcntl() with F_SETLK to release the
lock.
- This does not block; it just removes the
lock.
```

-----P_17-----

```
----- C Programming :
-----
```

```
===== init_ticket.c
=====
```

```
write(fd, &ticket, sizeof(ticket)) == -1:
->Writes the value of ticket (which is 0) into the file.
    ->&ticket ?
address of the integer variable.
    ->sizeof(ticket) ? number of bytes to write
(usually 4 bytes for int).
    ->If write() fails, prints error, closes the file,
and exits.
```

```
===== reserve_ticket.c =====
```

```
// Setup write lock
lock.l_type = F_WRLCK;
lock.l_whence = SEEK_SET;

lock.l_start = 0;
lock.l_len = 0; // Lock entire file
Configures a write lock:

short l_type;    // Type of lock: F_RDLCK, F_WRLCK, F_UNLCK
short l_whence;  // How
to interpret l_start (SEEK_SET (beginning) , SEEK_CUR, SEEK_END)
off_t l_start;   //
Starting offset for lock
off_t l_len;     // Number of bytes to lock (0 means till EOF)
```

```

pid_t l_pid;        // PID of process holding lock (for F_GETLK)

fcntl(fd,
F_SETLKW, &lock) == -1

lseek(fd, 0, SEEK_SET);
    -> Moves the file
pointer to the beginning of the file.

read(fd, &ticket, sizeof(ticket)) == -1 -->
Read current ticket number.
ticket++; --> inc ticket num

write(fd, &ticket,
sizeof(ticket)) == -1 --> write incremented ticket number.

lock.l_type = F_UNLCK;

fcntl(fd, F_SETLK,
&lock);

```

-----P_18-----

----- C Programming : -----

```

int lock_record(int fd, int recno, short lock_type)
    -fd -
an open file descriptor for the file you want to lock a record inside.
    -recno - the
record number (0,1,2,...) you want to lock.
    -lock_type - F_RDLCK (shared/read lock)
or F_WRLCK (exclusive/write lock).
    -Returns 0 on success, -1 on failure (and errno will
be set).

{
    struct flock lock;           // file lock variable
    lock.l_type =
lock_type;    // F_RDLCK or F_WRLCK or F_UNLCK
    lock.l_whence = SEEK_SET;    //
l_start is offset from file start.
    lock.l_start = recno * sizeof(int); // offset to
record
    lock.l_len = sizeof(int);    // lock only this record
    lock.l_pid
= getpid();

    fcntl(fd, F_SETLKW, &lock) == -1
}

int unlock_record(int fd,
int recno)
    -fd - an open file descriptor for the file you want to lock a record
inside.
    -recno - the record number (0,1,2,...) you want to lock.
{
    struct
flock lock;
    lock.l_type = F_UNLCK;
    lock.l_whence = SEEK_SET;

    lock.l_start = recno * sizeof(int);    // offset to record
    lock.l_len = sizeof(int);

    lock.l_pid = getpid();

    fcntl(fd, F_SETLK, &lock) == -1
}

off_t size = lseek(fd, 0, SEEK_END);
off_t lseek(int fd, off_t offset, int whence);

```

-> fd ? File descriptor (from open()).
 -> offset ? A number (how much
 to move).
 -> whence ? Where to start measuring from. Common values:
 -> SEEK_SET ? from the beginning of file.
 -> SEEK_CUR ? from the current
 position.
 -> SEEK_END ? from the end of
 file.

-----P_19-----

----- C Programming : -----

```
static __inline__ uint64_t rdtsc(void)
    -> static
? restricts visibility to the translation unit (source file).
    -> __inline__
? suggests compiler to inline the function (replace call with code).
    -> Return
type: uint64_t ? 64-bit unsigned integer.

__asm__ __volatile__("rdtsc" :
    "=a"(lo), "=d"(hi));
    -> __asm__ __volatile__ ? inline
assembly that the compiler should not optimize away or reorder.
    ->
"rdtsc" ? the CPU instruction Read Time Stamp Counter.
    -> It reads
the number of cycles since the last reset of the CPU.
    -> The result is a 64-bit
value:
    -> Lower 32 bits ? stored in register EAX (=a constraint ?
variable lo).
    -> Upper 32 bits ? stored in register EDX (=d constraint ?
variable hi).

Combining High and Low
    return (uint64_t)hi << 32 |
lo;
    -> (uint64_t)hi << 32 ? shift the high part left by 32 bits.

    -> | lo ? OR with the low part.
    -> Together ? produce
full 64-bit timestamp.

USE of RDTSC :
    -> rdtsc gives the
raw CPU cycle counter.
    -> Common uses:
    ->
High-resolution performance measurement (benchmarking).
    -> Profiling
code execution time at the CPU cycle level.
    -> Implementing timers or
random number seeds.

start = rdtsc(); -> Read timestamp before

pid =
getpid(); -> Call getpid()

end = rdtsc(); -> Read timestamp after

----- Get Freq -----

struct timespec ts_start, ts_end; // ts_start / ts_end: store wall-clock time before and after
a 1-second sleep.
    uint64_t start, end; // start / end: store the CPU cycle
count from rdtsc() at those times.
```

```

    clock_gettime(CLOCK_MONOTONIC, &ts_start); //
gets a monotonic timestamp (wall-clock time that never goes backward, not affected by system
clock changes).
    start = rdtsc(); // read the time-stamp
counter (number of CPU cycles since reset).

    sleep(1); // wait 1 second

clock_gettime(CLOCK_MONOTONIC, &ts_end); // gets a monotonic timestamp (wall-clock time
that never goes backward, not affected by system clock changes).
    end = rdtsc();
    // read the time-stamp counter (number of CPU cycles since reset).

// Add them to get total elapsed time as a double (in seconds).
    double elapsed_sec =
(ts_end.tv_sec - ts_start.tv_sec) +
    (ts_end.tv_nsec -
ts_start.tv_nsec) / 1e9; // Convert nanoseconds ? seconds (/ 1e9).
    return (end -
start) / elapsed_sec; // Hz (cycles per
second)

```

-----P_20-----

----- Commands: -----

./prio &
=>The & at the end:
-> In
Linux/Unix shells, appending & tells the shell to run the command in the background.

-> This means:
-> The process starts executing, but your terminal
does not wait for it to finish.
-> You immediately get back your shell
prompt so you can type more commands.

=>ps command:

->Stands for process status.
->Shows information about running processes.

->Without options, ps just lists processes belonging to the current shell.
->With options, you can customize what info to show.

-o pid,pri,ni,comm:

->-o flag lets you specify which columns/fields to display.
->

pid : Process ID

-> pri : priority value

-> ni : Nice

value (user-space priority hint).

-> Range: -20 (highest priority) to 19
(lowest priority). Default is 0.

-> comm : command name that started the
process

->Example: prio.

=> nice command:

->

Used to start a program with a given nice value (process priority hint).

->

Syntax:

-> nice -n <value> command [args...]

-> If no

-n is given, it defaults to 10.

\$nice -n 10 ./prio

=> -n 10

? sets nice value to +10.

=>./prio ? runs the executable prio from the

current directory.

=> What is "nice value"?

->The nice value (NI) is a user-space concept that influences how much CPU time the scheduler gives to the process.

->Range: -20 ? highest priority (needs root privileges). 0 ? default priority. +19 ? lowest priority.

=>

Linux kernel computes an internal priority (PRI) based on base priority + nice value.

----- C Programming : -----
infinite loop

-----P_21-----

----- C Programming : -----

fork() :
-> It is a system call that creates a new process.
-> After this line, two processes exist:
 Parent process ? The original process.
 Child process ? A duplicate of the parent.
-> Both processes resume execution from the point of the fork() call.
-> The return value of fork() is how each process distinguishes itself:
 In the parent process, fork() returns the child's PID (positive integer).
 In the child process, fork() returns 0.
 If an error occurs, it returns -1 in the parent (no child is created).

getpid() ? returns the process ID of the child itself.
getppid() ? returns the parent's process ID.

-----P_22-----

----- C Programming : -----

fd = open("output.txt", O_WRONLY | O_CREAT | O_APPEND, 0644);
-> Open file (create if it doesn't exist, write-only, append)

pid = fork();
-> Fork a child process

pid == 0? write(fd, child_text, strlen(child_text)) : write(fd, parent_text, strlen(parent_text));

-----P_23-----

----- C Programming : -----

pid = fork(); // Create a child process
// Child process
printf("Child process (PID = %d) is exiting...\n", getpid());

```

    // Parent process
    printf("Parent process (PID = %d) sleeping...\n",
getpid());
    printf("Child PID = %d will become a zombie until parent calls
wait()\n", pid);
    // Sleep to keep parent alive and child in zombie state

sleep(30);

```

----- Commands: -----

ps aux :

->This is a very common way to list processes in Linux:

->a ? show processes for all users (not just you).

->u ? show

processes in a user-oriented format (includes username, CPU%, MEM%, etc.).

->x ? include processes not attached to a terminal (daemons, background services).

->So ps aux lists all processes on the system, with lots of useful info.

| (pipe): [Important] [i/p Sequence of commands]

=> Sends the output of

the first command (ps aux) as input to the next command (grep Z).

grep Z :

->grep searches for lines containing Z.

? What is a Zombie Process?

-> A zombie

process is a process that has finished execution (it has exited), but still has an entry in the process table.

-> It remains because its parent process hasn't collected its exit status yet using wait() or waitpid().

? In ps output, its state appears as Z (Zombie) and often as <defunct> in the command column.

? Life Cycle of a Normal Process

->

Parent creates a child using fork().

-> Child executes some code (maybe runs another program via exec).

-> Child finishes and calls exit().

-> The kernel saves the child's exit status in the process table.

-> Parent calls wait() or waitpid().

-> Kernel gives the exit status to the parent and then removes the child's entry ? child is fully gone.

? Life Cycle of a Zombie

-> Parent creates a child with fork().

-> Child finishes execution and calls exit().

-> Kernel marks it as zombie and keeps its entry in the process table (to store exit code).

-> Parent does not call wait() (maybe by mistake or bad programming).

-> The child stays as a zombie (status Z) forever until parent dies or finally collects the status.

? Why Kernel Keeps Zombies?

-> The exit status contains info like:

-> Exit code (success/failure).

->

Resources used (CPU time, etc.).

-> Parent needs this info (via wait()), so the kernel cannot throw it away immediately.

? Problems with Zombies

->A single zombie is harmless (it takes no CPU, just an entry in process table).

->Many zombies = process

table fills up = new processes cannot be created.

->Indicates buggy parent code (not reaping children properly).

? Summary:

-> A zombie process = "dead but not cleaned up".

-> It's created when a child exits, but the parent doesn't read its exit status.

-----P_24-----

----- C Programming : -----

```
pid = fork(); // Create a child process

// Child process
printf("Child process
(PID = %d) running. Parent PID = %d\n", getpid(), getppid());
sleep(10); // Sleep so
parent exits first
printf("Child process (PID = %d) after parent exit. New Parent PID
= %d\n", getpid(), getppid());

// Parent process
printf("Parent process (PID
= %d) exiting...\n", getpid());
exit(0); // Parent exits immediately
```

Perfect

?, let's clear up Orphan processes and compare them with Zombie processes.

? What

is an Orphan Process?

-> An orphan process is a process whose parent has terminated (died) while the process itself is still running.

-> In Linux/Unix, when a parent dies, the orphan process is automatically adopted by init (PID 1, or nowadays systemd).

-> init becomes the new parent and is responsible for cleaning up when the orphan finishes.

? Life Cycle of an Orphan

-> A parent creates a child (via fork()).

-> Parent terminates without waiting for the child.

-> The child is still running

? becomes an orphan.

-> init (PID 1) adopts the orphan.

-> When the

orphan finishes, init calls wait() ? no zombies left behind.

Aspect	Zombie	Orphan
Definition	Child finished, parent didn't wait.	Child still running, parent finished.
State in ps	'Z' (Zombie, <defunct>).	'S', 'R', etc. (normal running/sleeping).
Parent	Still alive but ignoring child. Dead, replaced by init.	
Cleanup	Parent must wait(), else entry stays.	init reaps when child ends ? no permanent issue.
Problem	Can accumulate, filling process table.	Harmless, system handles automatically.

C Programming :

```
pid1 = fork();
pid2 = fork();
pid3 = fork();
```

? Step 1:

```
wait() / waitpid() :
    waited_pid = waitpid(pid2, &status, 0);
    pid_t waited_pid =
waitpid(child_pid, &status, 0);
    waitpid() returns:
        >0 ? PID of the
terminated child.
        0 ? no child has exited yet (only in non-blocking mode).
```

-1 ? error (e.g., no children).

```
if (waited_pid > 0) {
    if
(WIFEXITED(status)) {
        printf("Parent: Child 2 exited with status %d\n",
WEXITSTATUS(status));
    }
}
```

? Step 2: status

-> status is an integer that
encodes information about how the child ended.
-> We don't read it directly;
instead, we use macros like WIFEXITED, WEXITSTATUS, WIFSIGNALED, etc.

? Step 3:

```
WIFEXITED(status)
    -> Returns true (nonzero) if the child terminated normally (via exit()
or by returning from main).
    -> Returns false if the child was killed by a signal (e.g.,
SIGKILL).
```

? Step 4: WEXITSTATUS(status)

-> If WIFEXITED(status) is true, then
WEXITSTATUS(status) gives the exit code passed by the child.
->Example:

```
->exit(5);
    ? Parent sees WEXITSTATUS(status) == 5.
```

? Summary:

-> waitpid()
tells the parent which child ended.
-> WIFEXITED(status) checks if child exited
normally.
-> WEXITSTATUS(status) extracts the exit code from the
child.

C Programming :

Program A: use some executable program

```
char *program =
"./a.out"; // Program to execute
char *argv[] = {program, NULL}; // Argument
list (terminated with NULL)
```

```
if(execv(program, argv) < 0) {
    perror("execv
```

```
failed");
    return 1;
}
```

-> `execv()` replaces the current process image with a new one (here `./a.out`).

=> On success:

-> The original program code is

discarded.

-> Memory is overwritten by the new program.

-> Execution

restarts from the new program's `main()`.

-> Importantly: `execv()` does not create a

new process. The PID stays the same, but the code and data are replaced.

=> On

failure:

-> Returns -1.

-> Sets `errno`.

-> `perror("execv`

failed") prints the error message.

-> Program exits with status 1.

=>

This line only executes if `execv()` fails, because on success, the current process is completely replaced by the new program image.

=> Normal termination if reached (which won't happen if `execv()` succeeds).

? Summary

=> This program demonstrates `execv()`

replacing the current process with another program.

=> Steps:

- Print

the PID and the target program.

- Call `execv("./a.out", argv)`.

- If it succeeds, the current process is replaced, so execution continues inside the new `./a.out`.

- If it fails, `perror` reports the error, and the program exits with code

1.

=> Key point: Unlike `fork()`, `execv()` does not create a new process. The process is the same, but its code and memory are replaced.

Program B: pass some input to an

executable program. (for example execute an executable of `./a.out` name)

```
char *program
```

```
= "./a.out"; // Executable to run
```

```
char *arg1 = "Shikhar";
```

```
// Input argument
```

```
char *argv[] = {program, arg1, NULL}; // Argument list
```

```
(NULL-terminated)
```

```
// Replace current process with executable
```

```
if(execv(program, argv)
```

```
< 0) {
```

```
    perror("execv failed");
```

```
    return 1;
```

```
}
```

-----P_27-----

----- C Programming : -----

P - 1 ;

=> `execl("/bin/ls", "ls",`

"-Rl", NULL);

`execl()` is one of the `exec` family of functions.

Syntax:

```
int execl(const char *path, const char *arg0, ..., NULL);
```

```
-> path ? full path of the program to run (/bin/ls).
    -> arg0 ?
traditionally the program name (ls).
    -> arg1, arg2, ... ? command-line
arguments (-Rl).
    -> NULL ? marks the end of the argument list.
```

? Here:

```
Path = /bin/ls (actual program file on disk).
First argument
(argv[0]) = "ls".
Second argument = "-Rl".
Equivalent shell
command =
```

```
=> What Happens Inside execl?
    execl() replaces the current
process image with the new program (/bin/ls).
    -> That means:
```

```
-> The code of your program (main()) is overwritten.
    -> From this point,
your program ceases to exist.
    -> The process now becomes /bin/ls.
```

```
-> It does not return if successful.
    So, after execl(), the process starts
executing ls -Rl inside the same process ID.
```

```
P - 2 :
=>char *envp[] = {
"PATH=/bin:/usr/bin", NULL };
    - Here you define a custom environment for the
new program.
    - Each string is of the form:
=> execl("/bin/ls",
"ls", "-Rl", NULL, envp);
    - execl() is like execl(), but with
extra environment parameter.
=> int execl(const char *path, const char *arg0, ...,
NULL, char * const envp[]);
    - path ? full path of the program (/bin/ls).
    -
arg0 ? program name (ls).
    - arg1 ? "-Rl".
    - NULL ? marks the
end of arguments.
    - envp ? environment variables for the new program.
```

P - 3 :

```
=> execlp("ls", "ls", "-Rl", NULL); // Uses PATH to find ls

    - Calls execlp() from the exec family. Prototype (simplified): int execlp(const char *file,
const char *arg0, ..., (char *)NULL);
    -> Behavior:
        -> execlp searches
for the executable named by the first parameter ("ls") using the PATH environment
variable (so you don't need to provide /bin/ls explicitly).
        -> The
remaining arguments form the new process's argv[] list. By convention argv[0] is the
program name - here also "ls".
        -> "-Rl" is passed as
argv[1] (it is equivalent to -R -l for ls).
        -> The argument list must be
terminated with NULL.
        -> On success: the current process image is replaced by
the ls program. Execution does not return to this program; ls runs in the same PID, inheriting
open file descriptors and environment.
        -> On failure: execlp returns -1 and
errno is set.
```

```

P - 4 :
=> char *argv[] = { "ls", "-Rl", NULL };
Entry point.
-> Defines an argument vector argv:
-> argv[0] =
"ls" ? by convention, the first element is the program name.
->
argv[1] = "-Rl" ? option to ls.
-> argv[2] = NULL ? terminates the
array (required).
=> execv("/bin/ls", argv); // argv array, first element =
program name
=> Calls execv() with:
-> The absolute path /bin/ls ?
no PATH lookup (unlike execlp).
-> The argument vector argv.

=>
Behavior:
-> If successful, the current process image is replaced by /bin/ls.

-> PID stays the same, but code, data, and stack are replaced.
->
Execution does not return to this program – instead, it starts running the new program's
main().
? Summary:
-> This program prints a message, then replaces itself with
/bin/ls -Rl. If successful, you'll only see the ls output – the perror and return 1 lines
won't execute.

P - 5 :

=> char *argv[] = { "ls", "-Rl", NULL };

Program entry point.
Defines the argument vector argv for the new program:

->argv[0] = "ls" ? by convention, the program name.

->argv[1] = "-Rl" ? options passed to ls (recursive, long format).

->argv[2] = NULL ? marks the end of the argument list.

=>
execvp("ls", argv); // Uses PATH to find ls
Calls execvp() to replace the
current process with a new one running ls.
=> Parameters:
"ls" ? the
filename to execute (does not need a full path).
argv ? the argument vector (same as
would be passed in main(int argc, char *argv[])).
=> execvp() behavior:
->
Uses the PATH environment variable to locate the ls executable (so it finds /bin/ls or
/usr/bin/ls).
-> If successful:
-> The current process image is
replaced by ls.
-> The PID remains the same.
-> Control never
returns to this program – instead, ls starts running immediately.
-> If it fails
(e.g., "ls" not found in PATH), it returns
-1.

```

```

max_fifo = sched_get_priority_max(SCHED_FIFO);
-> min_fifo =
sched_get_priority_min(SCHED_FIFO);

```

sched_get_priority_max(SCHED_FIFO) ? returns the highest priority value available for the SCHED_FIFO scheduling policy.

sched_get_priority_min(SCHED_FIFO) ? returns the lowest priority value available for the same policy.

? These functions are necessary because the numeric range of real-time priorities is platform-dependent. On Linux, typically:
 (but you shouldn't hard-code these values, hence the API calls).

=> Maximum and minimum priority for SCHED_RR

```

-> max_rr = sched_get_priority_max(SCHED_RR);
-> min_rr =
sched_get_priority_min(SCHED_RR);
    Same thing, but for the Round-Robin real-time
policy.

```

On Linux, SCHED_FIFO and SCHED_RR usually share the same priority range (1 to 99).

? Key concepts

SCHED_FIFO:

```

-> First-In, First-Out scheduling for
real-time tasks.
-> A task runs until it voluntarily yields, blocks, or is preempted
by a higher-priority task.
-> No time slicing between tasks of equal priority.

```

SCHED_RR:

```

-> Round-Robin scheduling for real-time tasks.
-> Similar to

```

SCHED_FIFO, but tasks of the same priority take turns (each gets a timeslice).
 Priorities:

```

-> Real-time scheduling policies use integer priorities.
-> Higher number
= higher priority.
-> Range (on Linux): 1-99.
-> Normal (time-sharing)

```

processes with SCHED_OTHER all use priority 0 (their "niceness" value is separate from real-time priorities).

? Summary:

-> This program queries the OS for the valid priority ranges for the two main real-time scheduling policies (SCHED_FIFO and SCHED_RR) and prints them.

-----P_29-----

----- C Programming :

```

pid_t pid = getpid(); // Current process ID
int policy;

struct sched_param param;
->sched_param ? a struct that holds scheduling
parameters. Its main field is sched_priority (the priority of the process under real-time
policies).

-> policy = sched_getscheduler(pid);
if(policy == -1) {

perror("sched_getscheduler failed");
return 1;

}

```

`sched_getscheduler(pid)` ? queries the current scheduling policy of the given process (here, `pid` = current process).

Returns one of:

-> `SCHED_OTHER` ? normal time-sharing

policy (default for most processes).

-> `SCHED_FIFO` ? first-in, first-out real-time

scheduling.

-> `SCHED_RR` ? round-robin real-time scheduling.

If it fails,
prints an error.

-> Set new scheduling policy to `SCHED_FIFO` with maximum priority

-> `param.sched_priority = sched_get_priority_max(SCHED_FIFO);`

->

`sched_get_priority_max(SCHED_FIFO)` ? returns the maximum priority allowed for the `SCHED_FIFO` policy (usually 99 on Linux).

-> Sets `param.sched_priority` to that value.

->

`if(sched_setscheduler(pid, SCHED_FIFO, ¶m) == -1) {`

`perror("sched_setscheduler failed");`

`return 1;`

`}`

->

`sched_setscheduler(pid, SCHED_FIFO, ¶m)` ? attempts to set the scheduling policy of this process (`pid`) to `SCHED_FIFO` with the given priority.

-> Requires root (superuser) privileges. If run as a normal user, you'll get:

? Key Takeaways

-> Default

Linux policy: Processes usually run under `SCHED_OTHER` (normal, time-sharing).

->

Real-time policies: `SCHED_FIFO` and `SCHED_RR` require root privileges because they can hog the CPU if misused.

-> Priorities:

`SCHED_OTHER` ignores `sched_priority`.

`SCHED_FIFO` / `SCHED_RR` use a range of 1-99 (higher = more urgent).

-> Effect: If

successful, your process will run with real-time priority, which means it could starve normal processes.

-----P_30-----

----- C Programming : -----

=> `umask(0);` // Set file permissions mask

-> `umask(0)`

removes any default file permission restrictions.

-> Daemon-created files will have

exactly the permissions requested.

=> `sid = setsid();`

-> `setsid()` creates a new

session and makes this process the session leader.

-> This detaches the process from the controlling terminal ? crucial for daemons.

-> Now it runs independently.

=>

`chdir("/");` // Commented out to allow relative paths

Normally, a daemon

changes working directory to `/` to avoid blocking file systems from being unmounted.

->

Here it's commented out so the daemon can still use relative paths.

```
=> Close standard
file descriptors
    close(STDIN_FILENO);
    close(STDOUT_FILENO);
```

```
close(STDERR_FILENO);
    -> A daemon doesn't need a terminal.
    ->
Closing these prevents accidental reads/writes to terminal.
```

=> Daemon loop:

```
while (1) {
    time_t now = time(NULL);
    struct tm *t = localtime(&now);

}
    -> Infinite loop: the daemon keeps running forever.
    ->
time(NULL) ? current time in seconds.
    -> localtime() ? convert to local time
(hours, minutes, etc.).

    // Set target time (for testing, a minute ahead)

    int target_hour = t->tm_hour;
    int target_minute = (t->tm_min + 1) %
60;
    -> For demonstration: it sets the target time to one minute ahead of
current time.
    -> target_hour = current hour.
    ->
target_minute = current minute + 1 (wraps around with % 60).

    if (t->tm_hour ==
target_hour && t->tm_min == target_minute) {
        // Run the script (absolute
path, no spaces)
        system("/home/shikhar/myscript.sh");

        // Avoid
multiple runs in the same minute
        sleep(60);
    }
    else {

sleep(10);
    }

    -> Checks if current time matches the scheduled time.

-> If yes ? runs the script /home/shikhar/myscript.sh.
    -> Then sleeps for 60
seconds so it doesn't run multiple times in the same minute.
    -> Otherwise, sleeps
for 10 seconds and checks again.

    ? Key Concepts
    Daemonization steps:

    -> Fork and exit parent.
        -> Create new session (setsid).

-> Optionally change directory to /.
    -> Set umask(0).
    -> Close
standard file descriptors.
    Daemon behavior:
        -> Runs in background
without terminal.
        -> Independent of user sessions.
        -> Keeps
looping, checking system time, and executing tasks.
    In this program:
        ->
```


It executes a script exactly one minute later, then every hour/minute condition can be extended.

-> Works like a simplified custom cron service.